


# Univariate Time Series Analysis

Kevin Sheppard

# Seasonality

- A seasonal time series has a deterministic pattern that repeats on an annual basis


In [5]: `plot(housing, y=-2)`



# Seasonal Autocorrelation Pattern

- Seasonal data has dynamics at the annual frequency

```
In [6]: acf_pacf_plot(housing, 24, size=-2)
```




# Seasonal Differencing

- Seasonal differencing differences using the seasonal period  $s$

$$\Delta_s Y_t = Y_t - Y_{t-s}$$

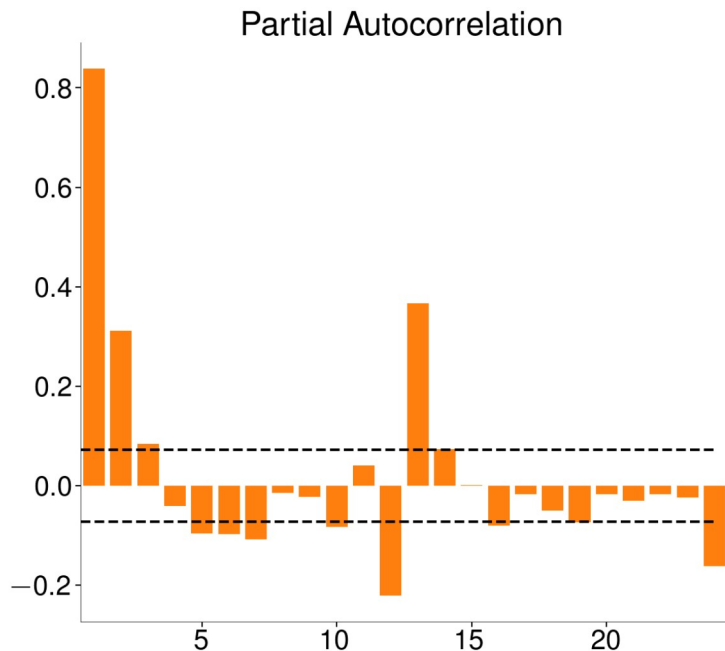
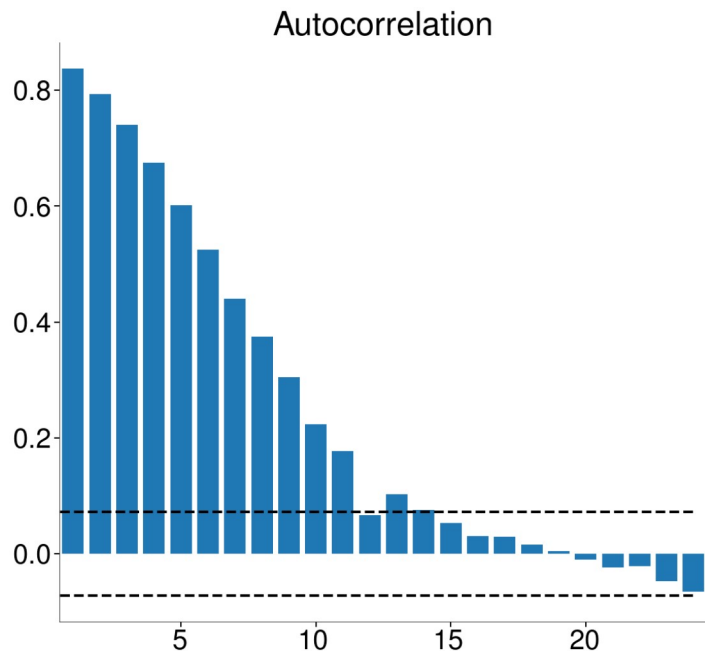
- Can reduce or eliminate seasonalities

```
In [7]: plot(housing_yoy, y=12)
```




# Seasonal Difference Autocorrelation Pattern

```
In [8]: acf_pacf_plot(housing_yoy, 24)
```



# Modeling Seasonal Differences

```
In [9]: res = SARIMAX(housing_yoy, order=(1, 0, 0)).fit()  
resids = res.resid.iloc[1:]  
acf_pacf_plot(resids, 24, size=-2)
```



# Seasonal ARMA Models

- SARMA models data at both the observational and seasonal frequency
- In lag polynomial representation

$$(1 - \phi_1 L - \phi_2 L^2)Y_t = \phi_0 + (1 + \theta_s L^s)\epsilon_t$$

- In standard ARMA representation


$$Y_t = \phi_0 + \phi_1 Y_{t-1} + \phi_2 Y_{t-2} + \theta_s \epsilon_{t-s} + \epsilon_t$$
$$\text{SARMA}(P, 0, Q) \times (P_s, 0, Q_s, s)$$

```
In [10]: res = SARIMAX(housing_yoy, order=(2, 0, 0), seasonal_order=(0, 0, 1, 12)).fit()  
summary(res)
```

	coef	std err	z	P> z	[0.025	0.975]
<b>ar.L1</b>	0.6809	0.034	20.284	0.000	0.615	0.747
<b>ar.L2</b>	0.2824	0.034	8.233	0.000	0.215	0.350
<b>ma.S.L12</b>	-0.8795	0.017	-50.520	0.000	-0.914	-0.845
<b>sigma2</b>	0.0083	0.000	21.791	0.000	0.008	0.009

# SARMA Residual Diagnostics

```
In [11]: acf_pacf_plot(res.resid.iloc[13:], 24)
```





# Extended Dynamics $(1, 0, 1) \times (1, 0, 1)_{12}$

- In lag polynomial representation

$$(1 - \phi_1 L)(1 - \phi_s L^s)Y_t = \phi_0 + (1 + \theta_1 L)(1 + \theta_s L^s)\epsilon_t$$

- In standard ARMA representation


$$Y_t = \phi_0 + \phi_1 Y_{t-1} + \phi_s Y_{t-s} - \phi_1 \phi_s Y_{t-s-1} + \theta_1 \epsilon_{t-1} + \theta_s \epsilon_{t-s} + \theta_1 \theta_s \epsilon_{t-s-1} + \epsilon_t$$

```
In [12]: res = SARIMAX(housing_raw, order=(1, 0, 1), seasonal_order=(1, 0, 1, 12)).fit()  
summary(res)
```

	coef	std err	z	P> z	[0.025	0.975]
<b>ar.L1</b>	0.9994	0.001	1441.218	0.000	0.998	1.001
<b>ma.L1</b>	-0.3196	0.031	-10.367	0.000	-0.380	-0.259
<b>ar.S.L12</b>	0.9987	0.001	1469.955	0.000	0.997	1.000
<b>ma.S.L12</b>	-0.8957	0.016	-54.402	0.000	-0.928	-0.863
<b>sigma2</b>	0.0083	0.000	22.277	0.000	0.008	0.009

# Extended Model Residual Diagnostics

```
In [13]: acf_pacf_plot(res.resid.iloc[13:], 24)
```



# Incorporating differencing

- SARIMA - Seasonal Autoregressive Integrated Moving Average

$$\text{SARIMA}(P, D, Q) \times (P_s, D_s, Q_s, s) \\ \Phi(L)\Phi_s(L)(1-L)^D(1-L^s)^{D_s}Y_t = \phi_0 + \Theta(L)\Theta_s(L)\epsilon_t$$

- **AR**

$$\Phi(L) = 1 - \phi_1 L - \phi_2 L^2 - \dots - \phi_P L^P$$

- **Seasonal AR:**

$$\Phi_s(L) = 1 - \phi_s L^s - \phi_{2s} L^{2s} - \dots - \phi_{P_s s} L^{P_s s}$$

- **MA:**

$$\Theta(L) = 1 + \theta_1 L + \theta_2 L^2 + \dots + \theta_Q L^Q$$

- **Seasonal MA:**

$$\Theta_s(L) = 1 + \theta_s L^s + \theta_{2s} L^{2s} + \dots + \theta_{Q_s s} L^{Q_s s}$$

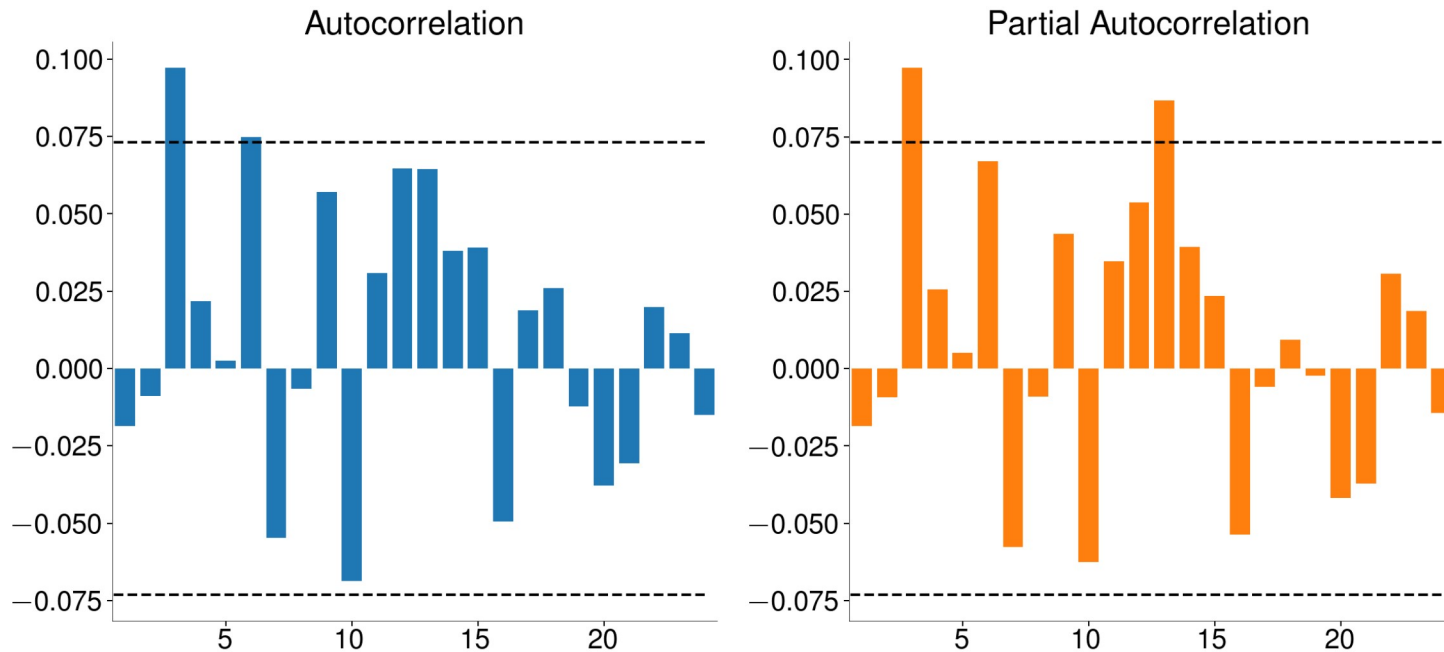
# Estimating SARIMA Models

```
In [14]: res = SARIMAX(housing_raw, order=(1, 0, 1), seasonal_order=(0, 1, 1, 12)).fit()  
summary(res)
```

	coef	std err	z	P> z	[0.025	0.975]
<b>ar.L1</b>	0.9777	0.008	126.932	0.000	0.963	0.993
<b>ma.L1</b>	-0.3128	0.033	-9.352	0.000	-0.378	-0.247
<b>ma.S.L12</b>	-0.8775	0.018	-48.061	0.000	-0.913	-0.842
<b>sigma2</b>	0.0083	0.000	21.771	0.000	0.008	0.009

# SARIMA(1, 0, 1) $\times$ (0, 1, 1)<sub>12</sub> Residual Analysis

```
In [15]: acf_pacf_plot(res.resid.iloc[13:], 24)
```



# Enforcing Unit Roots

- Set both differences to 1

- $D = 1$
- $D_s = 1$

- Model is


$$(1 - L)(1 - L^s)Y_t = \theta_s \epsilon_{t-s} + \epsilon_t$$
$$\Delta \Delta_s Y_t = \theta_s \epsilon_{t-s} + \epsilon_t$$

```
In [16]: res = SARIMAX(housing_raw, order=(0, 1, 1), seasonal_order=(0, 1, 1, 12)).fit()  
summary(res)
```

	coef	std err	z	P> z	[0.025	0.975]
<b>ma.L1</b>	-0.3246	0.031	-10.362	0.000	-0.386	-0.263
<b>ma.S.L12</b>	-0.8774	0.019	-47.035	0.000	-0.914	-0.841
<b>sigma2</b>	0.0084	0.000	22.025	0.000	0.008	0.009

# Double Difference Residual Diagnostics

```
In [17]: acf_pacf_plot(res.resid.iloc[13:], 24)
```



# Seasonal Dummies

- Alternative approach to SARIMA
- Model shifts as a deterministic processes
- One dummy for each of  $s - 1$  periods
  - Avoid dummy variable trap

```
In [18]: month = housing_raw.index.strftime("%b")
dummies = pd.get_dummies(month)
dummies = dummies[
    ["Feb", "Mar", "Apr", "May", "Jun", "Jul", "Aug", "Sep", "Oct", "Nov", "Dec"]
]
dummies.index = housing_raw.index
dummies.head()
```

Out[18]:

	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
1960-01-01	0	0	0	0	0	0	0	0	0	0	0
1960-02-01	1	0	0	0	0	0	0	0	0	0	0
1960-03-01	0	1	0	0	0	0	0	0	0	0	0
1960-04-01	0	0	1	0	0	0	0	0	0	0	0
1960-05-01	0	0	0	1	0	0	0	0	0	0	0




# Incorporating Dummies

```
In [19]: res = SARIMAX(housing_raw, order=(2, 1, 0), exog=dummies, trend="c").fit()  
summary(res)
```

	<b>coef</b>	<b>std err</b>	<b>z</b>	<b>P&gt; z </b>	<b>[0.025</b>	<b>0.975]</b>
<b>intercept</b>	0.0002	0.004	0.048	0.962	-0.007	0.008
<b>Feb</b>	0.0358	0.012	2.965	0.003	0.012	0.059
<b>Mar</b>	0.3075	0.012	24.777	0.000	0.283	0.332
<b>Apr</b>	0.4289	0.015	29.516	0.000	0.400	0.457
<b>May</b>	0.4669	0.018	26.260	0.000	0.432	0.502
<b>Jun</b>	0.4697	0.019	25.308	0.000	0.433	0.506
<b>Jul</b>	0.4328	0.019	23.265	0.000	0.396	0.469
<b>Aug</b>	0.4117	0.019	22.227	0.000	0.375	0.448
<b>Sep</b>	0.3657	0.017	21.802	0.000	0.333	0.399
<b>Oct</b>	0.3921	0.015	26.252	0.000	0.363	0.421
<b>Nov</b>	0.2169	0.013	16.942	0.000	0.192	0.242
<b>Dec</b>	0.0502	0.010	5.240	0.000	0.031	0.069
<b>ar.L1</b>	-0.2675	0.033	-8.115	0.000	-0.332	-0.203
<b>ar.L2</b>	-0.1107	0.034	-3.276	0.001	-0.177	-0.044
<b>sigma2</b>	0.0087	0.000	21.344	0.000	0.008	0.010

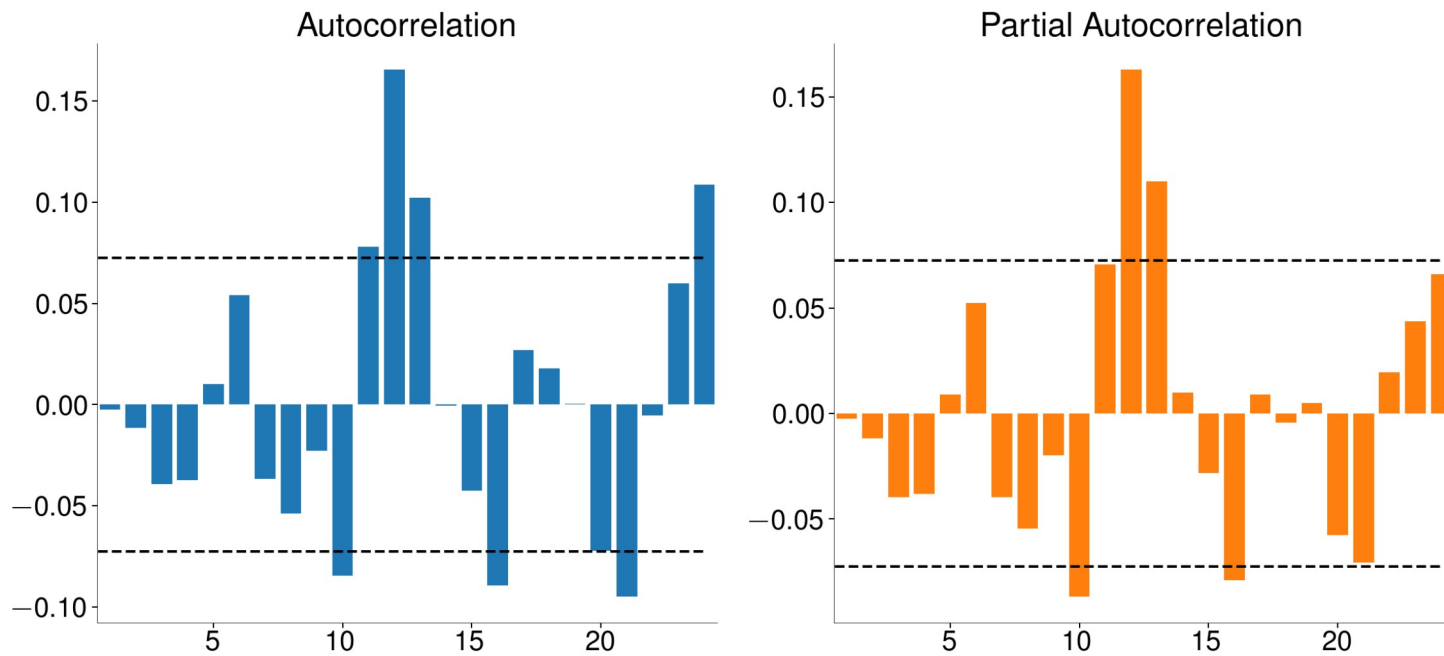
# Seasonal Dummy Model Residuals

```
In [20]: resid = res.resid.iloc[2:]  
plot(resid, y=-2)
```



# Seasonal Dummy Model Residuals Diagnostics

```
In [21]: acf_pacf_plot(resid, 24)
```



# Unit Root Testing


- Test  $Y_t$  for a unit root
- If null is not rejected, ensure that all required deterministic terms are included
- If null is still *not* rejected, difference and test  $\Delta Y_t$
- Positive  $Y_t$  series should usually be transformed with  $\ln$
- Select lag length of  $\Delta Y_{t-j}$  terms using IC, usually AIC
- $H_0 : \gamma = 0, H_1 : \gamma < 0$

$$\Delta Y_t = \delta_0 + \delta_1 t + \gamma Y_{t-1} + \sum_{i=1}^P \Delta \lambda_i Y_{t-i} + \epsilon_t$$

- Deterministic terms are needed to remove deterministic components from  $Y_{t-1}$

# The Dickey-Fuller distributions

```
In [23]: adf_cv_plot()
```



# Testing Default Premium

```
In [24]: from arch.unitroot import ADF
```

```
adf = ADF(default)  
adf.summary()
```

Out[24]: Augmented Dickey-

Fuller Results

Test Statistic	-3.866
P-value	0.002
Lags	10

Trend: Constant

Critical Values: -3.44 (1%), -2.87 (5%), -2.57 (10%)

Null Hypothesis: The process contains a unit root.

Alternative Hypothesis: The process is weakly stationary.

# Testing Curvature

```
In [25]: adf = ADF(curve)
adf.summary()
```

Out[25]: Augmented Dickey-  
Fuller Results

Test Statistic	-4.412
P-value	0.000
Lags	19

Trend: Constant


Critical Values: -3.44 (1%), -2.87 (5%), -2.57 (10%)

Null Hypothesis: The process contains a unit root.

Alternative Hypothesis: The process is weakly stationary.

# Testing Industrial Production

```
In [27]: plot(ln_indpro)
```





# ADF with a constant

```
In [28]: adf = ADF(ln_indpro, trend="c")  
adf.summary()
```

Out[28]: Augmented Dickey-  
Fuller Results

Test Statistic	-2.186
P-value	0.211
Lags	4

Trend: Constant

Critical Values: -3.44 (1%), -2.87 (5%), -2.57 (10%)

Null Hypothesis: The process contains a unit root.

Alternative Hypothesis: The process is weakly stationary.

# ADF Regression Results

```
In [29]: summary(adf.regression)
```

	coef	std err	t	P> t	[0.025	0.975]
<b>Level.L1</b>	-0.0017	0.001	-2.186	0.029	-0.003	-0.000
<b>Diff.L1</b>	0.3637	0.037	9.860	0.000	0.291	0.436
<b>Diff.L2</b>	-0.1110	0.039	-2.832	0.005	-0.188	-0.034
<b>Diff.L3</b>	0.0447	0.039	1.138	0.255	-0.032	0.122
<b>Diff.L4</b>	0.0217	0.037	0.589	0.556	-0.051	0.094
<b>const</b>	0.0083	0.003	2.579	0.010	0.002	0.015

# Increasing the deterministic order

```
In [30]: adf = ADF(ln_indpro, trend="ct")  
adf
```

Out[30]: Augmented Dickey-  
Fuller Results

Test Statistic	-1.831
P-value	0.690
Lags	6

Trend: Constant and Linear Time Trend

Critical Values: -3.97 (1%), -3.42 (5%), -3.13 (10%)

Null Hypothesis: The process contains a unit root.

Alternative Hypothesis: The process is weakly stationary.

# ADF Regression Results

```
In [31]: summary(adf.regression)
```

	<b>coef</b>	<b>std err</b>	<b>t</b>	<b>P&gt; t </b>	<b>[0.025</b>	<b>0.975]</b>
<b>Level.L1</b>	-0.0061	0.003	-1.831	0.067	-0.013	0.000
<b>Diff.L1</b>	0.3680	0.037	10.010	0.000	0.296	0.440
<b>Diff.L2</b>	-0.1034	0.039	-2.643	0.008	-0.180	-0.027
<b>Diff.L3</b>	0.0459	0.039	1.168	0.243	-0.031	0.123
<b>Diff.L4</b>	0.0489	0.039	1.245	0.213	-0.028	0.126
<b>Diff.L5</b>	-0.0571	0.039	-1.455	0.146	-0.134	0.020
<b>Diff.L6</b>	0.0646	0.038	1.701	0.089	-0.010	0.139
<b>const</b>	0.0231	0.011	2.053	0.040	0.001	0.045
<b>trend</b>	9.34e-06	7.07e-06	1.320	0.187	-4.55e-06	2.32e-05

# Quadratic Time Trend

$$\Delta Y_t = \delta_0 + \delta_1 t + \delta_2 t^2 + \gamma Y_{t-1} + \sum_{i=1}^P \lambda_i \Delta Y_{t-i} + \epsilon_t$$

```
In [32]: adf = ADF(ln_indpro, trend="ctt")  
adf
```

Out[32]: Augmented Dickey-  
Fuller Results

Test Statistic	-2.962
P-value	0.314
Lags	6

Trend: Constant, Linear and Quadratic Time Trends

Critical Values: -4.39 (1%), -3.84 (5%), -3.56 (10%)

Null Hypothesis: The process contains a unit root.

Alternative Hypothesis: The process is weakly stationary.

# ADF Regression Results

```
In [33]: summary(adf.regression)
```

	<b>coef</b>	<b>std err</b>	<b>t</b>	<b>P&gt; t </b>	<b>[0.025</b>	<b>0.975]</b>
<b>Level.L1</b>	-0.0157	0.005	-2.962	0.003	-0.026	-0.005
<b>Diff.L1</b>	0.3705	0.037	10.103	0.000	0.298	0.442
<b>Diff.L2</b>	-0.0987	0.039	-2.527	0.012	-0.175	-0.022
<b>Diff.L3</b>	0.0501	0.039	1.277	0.202	-0.027	0.127
<b>Diff.L4</b>	0.0538	0.039	1.372	0.170	-0.023	0.131
<b>Diff.L5</b>	-0.0526	0.039	-1.342	0.180	-0.129	0.024
<b>Diff.L6</b>	0.0713	0.038	1.878	0.061	-0.003	0.146
<b>const</b>	0.0526	0.017	3.105	0.002	0.019	0.086
<b>trend</b>	5.205e-05	1.97e-05	2.644	0.008	1.34e-05	9.07e-05
<b>quadratic_trend</b>	-3.109e-08	1.34e-08	-2.324	0.020	-5.73e-08	-4.82e-09

# Testing the difference

- The Industrial Productivity Index is  $I(1)$  since 1 difference makes it stationary

```
In [34]: delta_indpro = np.log(orig.INDPRO).diff().dropna()
adf = ADF(delta_indpro, trend="c")
adf
```

Out[34]: Augmented Dickey-Fuller

Results

Test Statistic	-11.945
P-value	0.000
Lags	3

Trend: Constant

Critical Values: -3.44 (1%), -2.87 (5%), -2.57 (10%)


Null Hypothesis: The process contains a unit root.

Alternative Hypothesis: The process is weakly stationary.

# The importance of time trends

- Incorrect trend specification results in no power to reject false nulls
- Always find a unit root in trending time series

```
In [36]: plot(trending, 13)
```





# Under-specifying Deterministic Terms

- Do not use model with no deterministic terms on real data unless  $Y_t$  is known to be mean 0

```
In [37]: ADF(trending, trend="n")
```

```
Out[37]: Augmented Dickey-
```

Fuller Results

Test Statistic	1.934
P-value	0.988
Lags	9

Trend: No Trend

Critical Values: -2.58 (1%), -1.94 (5%), -1.62 (10%)

Null Hypothesis: The process contains a unit root.

Alternative Hypothesis: The process is weakly stationary.

# Under-specifying Deterministic Terms

```
In [38]: ADF(trending, trend="c")
```

Out[38]: Augmented Dickey-Fuller Results

Test Statistic	-1.146
P-value	0.696
Lags	9

Trend: Constant

Critical Values: -3.46 (1%), -2.87 (5%), -2.57 (10%)

Null Hypothesis: The process contains a unit root.

Alternative Hypothesis: The process is weakly stationary.

# Correctly specifying Deterministic Terms

```
In [39]: ADF(trending, trend="ct")
```

Out[39]: Augmented Dickey-Fuller Results

Test Statistic	-6.790
P-value	0.000
Lags	0

Trend: Constant and Linear Time Trend

Critical Values: -4.00 (1%), -3.43 (5%), -3.14 (10%)

Null Hypothesis: The process contains a unit root.

Alternative Hypothesis: The process is weakly stationary.

# Over-specifying Deterministic Terms

```
In [40]: ADF(y, trend="ctt")
```

Out[40]: Augmented Dickey-Fuller Results

Test Statistic	-6.885
P-value	0.000
Lags	0

Trend: Constant, Linear and Quadratic Time Trends

Critical Values: -4.42 (1%), -3.86 (5%), -3.57 (10%)

Null Hypothesis: The process contains a unit root.

Alternative Hypothesis: The process is weakly stationary.

# Incorporating Time Trends into ARMA Models

- Can jointly estimate trends and stationary components of ARMA models


$$Y_t = \phi_0 + \delta_1 t + \phi_1 Y_{t-1} + \epsilon_t$$

```
In [41]: res = SARIMAX(trending, order=(1, 0, 0), trend="ct").fit()  
summary(res)
```

	coef	std err	z	P> z	[0.025	0.975]
<b>intercept</b>	0.0796	0.116	0.687	0.492	-0.148	0.307
<b>drift</b>	0.0255	0.004	6.149	0.000	0.017	0.034
<b>ar.L1</b>	0.6822	0.050	13.534	0.000	0.583	0.781
<b>sigma2</b>	0.9213	0.080	11.451	0.000	0.764	1.079

# Time Trend Model Residuals

```
In [42]: plot(res.resid.iloc[1:])
```



# Next Week

- Univariate Volatility Modeling